Batch No: _____

(Office Use)

CUSTOMER CHURN ANALYSIS
AND BUSINESS INTELLIGENCE DASHBOARD

A Summer Industry Internship Report
Submitted in Partial Fulfillment of the Requirements
For the award of the Degree of

Bachelor of Technology in

Computer Science and Engineering - Data Science (CSE-DS)

By

GAVUJI MYTHILI	**24311A67H8**
KONDAMADUGU ASHIWITH	**24311A67H6**
DOMA NIYATI	**24311A67H7**



DEPARTMENT OF DATA SCIENCE (CSE-DS)

SREENIDHI INSTITUTE OF SCIENCE AND TECHNOLOGY

November 2026

ACCEPTANCE BY INTERNAL TECHNICAL REVIEW COMMITTEE

(For Internal Department use only)

We hereby approve the Presentation and Report of the Summer Industry Internship work titled:

CUSTOMER CHURN ANALYSIS AND BUSINESS INTELLIGENCE DASHBOARD

Submitted by Roll No. **24311A67H8, 24311A67H6, 24311A67H7**, studying in the B.Tech **CSE-DATA SCIENCE** branch in **III rd Year I st Semester**.

Committee Members	Signature with Date
1. Name: Dr. Md. Jaffar Sadiq **Designation: Prof. & HoD**	_____
2. Name: Mr. Syed. Mazharuddin **Designation: Coordinator**	_____

****

For Faculty/ Evaluator Use Only:

If you choose not to sign the acceptance certificate above, please indicate the reason(s):

1. The amount of time and effort put in by the student is not sufficient
2. The report doesn't represent the actual work done/ expected to be done
3. The presentation of the seminar or report by the student is not up to the mark
4. Any other objection (please elaborate):

_____

SREENIDHI INSTITUTE OF SCIENCE AND TECHNOLOGY
(AUTONOMOUS)
DEPARTMENT OF CSE - DATA SCIENCE (CSE-DS)

CERTIFICATE

This is to certify that the Summer Industry Internship work titled **'CUSTOMER CHURN ANALYSIS AND BUSINESS INTELLIGENCE DASHBOARD'** submitted by the Students **G Mythili, K Ashiwith and D Niyati** bearing Roll No. **24311A67H8, 24311A67H6 and 24311A67H7**, towards partial fulfillment for the award of Bachelors Degree in Computer Science and Engineering Data Science (CSE-DS) from Sreenidhi Institute of Science and Technology, Ghatkesar, Hyderabad, is a record of Bonafide work done by him/ her. The results embodied in the work are not submitted to any other University or Institute for the award of any other degree or diploma.

Mr. Syed. Mazharuddin
Coordinator

Dr. Md. Jaffar Sadiq
HoD, CSE-DS

External Examiner Signature

_____

DECLARATION

We, **G Mythili (24311A67H8)**, **K Ashiwith (24311A67H6)**, and **D Niyati (24311A67H7)** students of Sreenidhi Institute of Science and Technology, Yamnampet, Ghatkesar, studying **III Year, I Semester**, Bachelor of Technology in **Computer Science and Engineering - Data Science (CSE-DS)**, hereby solemnly declare that the project report titled **'CUSTOMER CHURN ANALYSIS AND BUSINESS INTELLIGENCE DASHBOARD'** is submitted in partial fulfillment for the award of the degree of **Bachelor of Technology in Computer Science and Engineering - Data Science (CSE-DS)**.

It is further declared that, to the best of our knowledge, the work reported in this project has not been submitted earlier, either in part or in full, to any other University or Institute for the award of any degree or diploma.

G MYTHILI
24311A67H8 **SIGNATURE** _____

K ASHIWITH
24311A67H6 **SIGNATURE** _____

D NIYATI
24311A67H7 **SIGNATURE** _____

ACKNOWLEDGEMENT

It gives us immense satisfaction and profound pride to express our gratitude on the successful completion of our academic internship project. However, this undertaking would remain incomplete without acknowledging the support, guidance, and encouragement of the individuals who played an instrumental role in shaping this endeavor. We take this opportunity to sincerely thank all those who contributed directly and indirectly to the completion of this research project.

At the outset, we express our sincere thanks and deep sense of gratitude to **Mr. Syed. Mazharuddin**, Project Coordinator, for his continuous guidance, valuable suggestions, and constant encouragement throughout the course of this project. His insightful feedback, rigorous technical mentoring, and intellectual motivation helped us overcome intricate data engineering hurdles, formulate robust statistical hypotheses, and complete our work successfully.

We are extremely thankful to **Dr. Md. Jaffar Sadiq**, Head of the Department, CSE-DS, and our Principal, **Dr. T. Ch. Siva Reddy**, for providing us with the opportunity to undertake this project and for their consistent administrative support and encouragement. Their inspiring leadership and academic vision created an exemplary research environment that greatly contributed to the smooth execution of our project work.

Finally, we extend our sincere thanks to all faculty members of the Department of Data Science and our peers for their cooperation, technical suggestions, and constructive critique. Above all, we are forever grateful to our parents and families for their unconditional love, moral encouragement, and blessings, which formed the bedrock of our academic pursuits.

G MYTHILI	**24311A67H8**
K ASHIWITH	**24311A67H6**
D NIYATI	**24311A67H7**

ABSTRACT

Customer churn represents one of the most critical operational and financial vulnerabilities in the contemporary telecommunications industry, where acquiring a new customer costs up to five to seven times more than retaining an existing subscriber. This project presents an enterprise-grade Customer Churn Analysis and Business Intelligence Dashboard developed strictly utilizing rigorous Data Science, relational SQL database engineering, parametric and non-parametric statistical hypothesis testing, and exploratory data analytics without relying on black-box machine learning algorithms. The analytical pipeline ingests the comprehensive IBM Telco Customer Churn dataset comprising 7,043 customer accounts across 21 raw operational attributes. A multi-stage data quality remediation framework executes duplicate auditing, schema type-casting, whitespace anomaly resolution in total lifetime charges, and 1.5xIQR boundary outlier validation, establishing a pristine 28-dimensional analytical feature store.

Statistical evaluations encompass Chi-Square (Chi^2) tests of independence and Welch's two-sample independent t-tests to evaluate categorical and numerical defection dependencies. Empirical findings identify contract commitment as the single largest structural vulnerability, with Month-to-month contracts exhibiting a catastrophic 42.71% attrition rate generating 88.55% of all defection volume, compared to just 2.83% on Two-year agreements - a fifteen-fold stability disparity (Cramer's $V = 0.410$, $p < 0.0001$). Furthermore, unautomated payment friction via Electronic Check drives an alarming 45.29% churn rate, while customer defection is heavily front-loaded in the first twelve months (mean tenure of 17.98 months for churners vs 37.57 months for retained accounts, $p < 0.0001$). An optimized SQLite database engine executing ten business queries powers an interactive, production-ready Streamlit dashboard featuring reactive multi-cohort filtering, tenure-spend vulnerability matrices, and financial ROI simulation models that forecast over $328,000 in annual recurring revenue preservation through targeted term migration.

Keywords: Customer Churn Analysis, Business Intelligence, Exploratory Data Analysis, Relational Database Engineering, Chi-Square Test of Independence, Welch's T-Test, Telecom Retention Strategy, Streamlit Dashboard, Financial ROI Modeling.

TABLE OF CONTENTS

Section / Chapter Title	Page No.
Project Acceptance Certificate by Project Review Committee (PRC)	ii
Certificate of Completion from College	iii
Declaration	iv
Acknowledgement	v
Abstract	vi
List of Figures	viii
List of Tables	viii
Chapter 1: Introduction	**1**
1.1 General Context & Problem Statement	1
1.2 Motivation & Industry Relevance	1
1.3 Challenges & Shortcomings of Naive Approaches	2
1.4 Core Objectives & Research Scope	2
1.5 Report Organization	2
Chapter 2: Literature Survey & Background Study	**3**
2.1 Telecom Churn Management Evolution	3
2.2 Comparative Review: Rule-Based vs Data-Driven Systems	3
2.3 Research Gaps & Identified Limitations	4
2.4 Justification for Exploratory Data Science & SQL Paradigm	4
Chapter 3: Data Ingestion & Data Quality Engineering	**5**
3.1 IBM Telco Dataset Schema Overview	5
3.2 Ingestion Architecture & Data Auditing	5
3.3 Whitespace Type-Casting & Zero-Tenure Resolution	6
3.4 Duplicate Record Verification	6
3.5 Outlier Diagnostics via 1.5xIQR Fences	7
3.6 Automated Cleaning Pipeline Implementation	8
Chapter 4: Exploratory Data Analysis & Visual Intelligence	**9**
4.1 Univariate Distribution Analysis	9
4.2 Bivariate Churn Drivers & Contract Dynamics	10
4.3 Service Portfolio & Technology Impact	11
4.4 Financial Charges & Dispersion Analysis	12
4.5 Demographic Vulnerabilities: Seniors, Partners & Dependents	13
4.6 Multivariate Correlation Analysis & Heatmap	14
Chapter 5: Parametric & Non-Parametric Statistical Inference	**15**
5.1 Parametric & Distributional Metrics	15
5.2 Chi-Square Independence Tests & Cramer's V	16
5.3 Welch's Two-Sample Independent T-Tests	17
Chapter 6: Relational Database Modeling & Enterprise SQL	**18**
6.1 SQLite Schema Architecture & Indexing	18
6.2 Business Query Execution (Q1 to Q5)	19
6.3 Business Query Execution (Q6 to Q10)	20
6.4 SQL Performance Optimization & Query Plans	21
Chapter 7: Customer Segmentation & High-Risk Profiling	**22**
7.1 Lifecycle Tenure Cohorts	22
7.2 2D Tenure vs Spend Vulnerability Grid	23
7.3 Critical Defection Danger Zone Profiling	24
Chapter 8: Business Intelligence Dashboard Architecture	**25**
8.1 End-to-End System Architecture	25
8.2 Reactive State Pipeline & UI Components	26
8.3 Custom SaaS Dark Design System & Popover Drawers	27
Chapter 9: Strategic Recommendations & Financial ROI	**28**
9.1 Strategic Pillars 1 & 2: Contract Migration & Auto-Pay	28
9.2 Strategic Pillars 3 & 4: Onboarding Care & Bundling	29
9.3 Financial ROI Simulation Model	30
9.4 90-Day Executive Retention Implementation Roadmap	31
Chapter 10: Conclusion & Future Scope	**32**
Appendix I: Complete SQL Analytics Query Catalog (Part 1: Q1-Q5)	**33**
Appendix I: Complete SQL Analytics Query Catalog (Part 2: Q6-Q10)	**34**
Appendix II: Mathematical Formulations & Statistical Derivations	**35**
Appendix III: Complete Cleaned Data Dictionary & Schema	**36**
Appendix IV: PowerPoint Presentation Slides (Part 1: Slides 1-6)	**37**
Appendix IV: PowerPoint Presentation Slides (Part 2: Slides 7-12)	**38**
Appendix V: Software Module Architecture & File Manifest	**39**
References	**40**

LIST OF FIGURES

Fig. No.	Figure Title & Description	Page
Fig. 4.1	Overall Customer Churn Distribution Pie / Donut Chart	9
Fig. 4.2	Monthly Recurring Charges ($) Distribution Histogram	9
Fig. 4.3	Customer Churn Rate by Contract Agreement Type	10
Fig. 4.4	Customer Defection by Billing Payment Method	10
Fig. 4.5	Churn Disparity Across Internet Delivery Technologies	11
Fig. 4.6	Protective Impact of Value-Added Service Add-ons	11
Fig. 4.7	Customer Subscription Tenure Dispersion Boxplot	12
Fig. 4.8	Tenure vs. Monthly Charges Bivariate Scatter Distribution	12
Fig. 4.9	Demographic Churn Disparities (Senior Citizens & Dependents)	13
Fig. 4.10	Multivariate Numerical Correlation Heatmap	14
Fig. 7.1	Churn Rate Across Lifecycle Tenure Cohort Segments	22
Fig. 8.1	End-to-End System Architecture & Data Flow	25

LIST OF TABLES

Table No.	Table Title & Description	Page
Table 3.1	1.5xIQR Outlier Diagnostic Fences for Numerical Attributes	7
Table 3.2	Automated Data Quality Cleansing & Feature Engineering Matrix	8
Table 5.1	Parametric Summary Statistics of Numerical Attributes	15
Table 5.2	Chi-Square (Chi^2) Tests of Independence & Cramer's V Coefficients	16
Table 5.3	Welch's Two-Sample Independent T-Test Results	17
Table 6.1	SQLite Relational Database Table Schema & B-Tree Indexes	18
Table 6.2	SQL Execution Plan (EXPLAIN QUERY PLAN) Latency Benchmarks	21
Table 7.1	2D Tenure Cohort vs. Monthly Spend Churn Rate Matrix (%)	23
Table 9.1	Financial ROI Simulation Model for Churn Reduction Rates	30
Table 9.2	90-Day Executive Retention Implementation Roadmap	31
Table A.1	Complete Software Module & Source Code Manifest	39

CHAPTER 1

INTRODUCTION

1.1 General Context & Topic Relevance

In the contemporary global telecommunications industry, customer churn - defined as the rate at which subscribers discontinue their service contracts with an operator - constitutes the single greatest strategic, operational, and financial vulnerability. Modern telecommunication services operates in a mature, hyper-saturated marketplace characterized by aggressive pricing battles, rapid technological standardization, and minimal friction for customer defection. Leading telecommunication providers experience annual defection rates ranging between 20% and 35%, eroding recurring revenue bases and inflicting severe bottom-line damage.

Customer churn is not merely a marketing annoyance; it represents a fundamental economic threat to enterprise enterprise sustainability. When an account defects, the operator loses not only the immediate monthly recurring revenue (MRR) but also the amortized recovery of substantial upfront Customer Acquisition Costs (CAC), which include optical fiber installation, cellular hardware subsidies, marketing commissions, and account onboarding provisioning.

1.2 Motivation & Industry Relevance

Decades of empirical corporate finance and marketing literature consistently demonstrate that retaining an existing telecommunication subscriber costs between one-fifth and one-seventh the capital expenditure required to acquire a replacement customer (Blattberg et al., 2008). Furthermore, retained subscribers demonstrate substantially higher Customer Lifetime Value (CLV), expanding their consumption of premium broadband bandwidth, bundled entertainment services, and digital security add-ons over time.

Consequently, telecommunications executives recognize that reducing annual churn by even a modest margin of 2% to 5% preserves hundreds of thousands to tens of millions of dollars in predictable annual recurring cash flows. However, designing effective retention interventions requires granular, transparent diagnostic intelligence: executives cannot fix churn without mathematically isolating *which* service configurations cause defection, *which* customer cohorts are most vulnerable, and *where* friction resides in billing lifecycles.

1.3 Challenges & Shortcomings of Naive Approaches

Historically, telecommunications organizations addressed defection through primitive rule-based heuristic systems or uninterpretable black-box algorithms. Heuristic systems rely on static threshold rules - such as flagging customers who register three customer service calls or exceed a specific billing threshold. These naive approaches suffer from three catastrophic failure modes:

• **Extremely High False Positive Rates:** Static rules routinely flag genuine, highly engaged customers who simply called for routine inquiries, wasting expensive human customer-care intervention hours.

• **Complete Inability to Detect Structural Interaction Effects:** Heuristics fail to recognize that defection is rarely caused by a single isolated variable, but rather by compound multi-factor vulnerabilities (e.g., Month-to-month contracts interacting with unautomated electronic check billing and high fiber optic charges).

• **The Black-Box Dilemma of Complex ML Models:** While modern deep neural networks and gradient boosting ensembles can generate numerical defection probability scores, they provide zero causal interpretability to commercial executives. Business stakeholders cannot determine *why* a customer is defecting or *what* operational remedy to execute.

1.4 Core Objectives & Scope of Study

This project rejects black-box opacity in favor of a rigorous, transparent Data Science, SQL engineering, and Business Intelligence framework. The core objectives of this study are:

• To design and execute an end-to-end data quality pipeline that audits, cleanses, and enriches 7,043 customer accounts across 21 raw dimensions into a pristine 28-dimensional analytical store.

• To perform exhaustive univariate, bivariate, and multivariate exploratory data analysis isolating critical churn drivers.

• To conduct rigorous parametric and non-parametric statistical hypothesis tests (Chi-Square tests of independence, Cramer's V, and Welch's two-sample independent t-tests) to confirm mathematical significance.

• To engineer an enterprise relational SQLite database modeled with B-tree indexes executing ten core business queries.

• To deploy an interactive, executive-grade Streamlit Business Intelligence dashboard equipped with real-time cohort filtering and financial ROI retention calculators.

1.5 Organization of the Report

This report is organized into ten logical chapters: Chapter 2 conducts an exhaustive literature review; Chapter 3 details data quality engineering; Chapter 4 presents exploratory visual analytics; Chapter 5 details statistical hypothesis testing; Chapter 6 covers relational database modeling; Chapter 7 constructs customer segmentation matrices; Chapter 8 details dashboard software architecture; Chapter 9 outlines actionable executive retention pillars and ROI projections; and Chapter 10 presents conclusions, limitations, and future research roadmaps.

CHAPTER 2

LITERATURE SURVEY & BACKGROUND STUDY

2.1 Evolution of Telecom Churn Management

Academic and commercial inquiry into customer relationship management (CRM) and customer defection has evolved across three distinct paradigms over the past three decades. In the early telecommunications era (1990s), churn management was predominantly reactive: operators only addressed defection after the subscriber contacted customer support to request service disconnection. Retention departments offered ad-hoc fee waivers, which proved economically disastrous as discounts were frequently granted to accounts that had already made irrevocable switching commitments.

In the second evolutionary wave (2000s to 2015), database marketing and data warehousing introduced descriptive statistics and RFM (Recency, Frequency, Monetary) segmentation. Pioneering studies by Berson et al. (2000) and Blattberg et al. (2008) demonstrated that customer tenure and billing consistency served as significant indicators of brand commitment. However, these early frameworks suffered from batch processing limitations, where data warehouses took weeks to reconcile billing cycles, causing interventions to lag weeks behind actual defection triggers.

The modern paradigm (2016-present) emphasizes real-time analytical business intelligence, where relational database engines and interactive visual cockpits empower operational managers to monitor customer cohorts continuously, isolating defection triggers at the exact intersection of contractual terms, payment channels, and service add-ons.

2.2 Comparative Review: Rule-Based Systems vs Data-Driven Analytics

A critical comparison between traditional rule-based monitoring and comprehensive exploratory data analytics highlights substantial differences in operational efficacy:

• **Static Rule Systems:** Rely on rigid, hand-crafted business heuristics established by managerial intuition. Empirical research by Dal Pozzolo et al. (2015) proved that rule-based systems generate excessive false discovery rates, failing to scale as customer product catalogs diversify.

• **Black-Box Machine Learning:** Supervised classifiers such as Random Forests (Breiman, 2001) and XGBoost (Chen & Guestrin, 2016) maximize predictive accuracy. However, as highlighted by Rudin (2019) in *Stop Explaining Black Box Machine Learning Models for High Stakes Decisions*, black-box scores fail to establish domain trust. A business executive cannot restructure billing terms or modify contract pricing based purely on an uninterpretable probability score without understanding underlying statistical associations.

• **Exploratory Data Science & SQL Analytics:** This paradigm provides complete mathematical transparency. By utilizing descriptive statistics, contingency table hypothesis testing, and direct relational database querying, decision-makers obtain verifiable, auditable business truths.

2.3 Identified Literature Gaps & Practical Limitations

A thorough survey of published literature reveals three fundamental gaps in existing telecommunications churn research:

1. Overemphasis on Algorithmic Metric Optimization vs. Practical Interpretability: The vast majority of published studies focus exclusively on maximizing ROC-AUC or F1-scores using complex ensembles. They rarely examine whether the features driving predictions can be influenced by commercial business operations. Predicting that a customer will defect with 89% probability is practically useless if the operational drivers remain completely hidden.

2. Neglect of Payment Channel Friction and Operational Billing Mechanisms: Existing literature frequently treats payment method as an incidental demographic factor. Few empirical studies investigate how manual, friction-heavy settlement mechanisms - such as non-automated electronic checks - exacerbate monthly billing salience and trigger defection.

3. Disconnect Between Statistical Output and Financial ROI Modeling: Most academic frameworks conclude with confusion matrices and accuracy tables, failing to translate statistical insights into financial balance-sheet language: Monthly Recurring Revenue (MRR) preservation, Customer Lifetime Value (CLV), and budget ROI.

2.4 Justification for Exploratory Data Science & SQL Paradigm

To bridge these critical academic and operational literature gaps, this project adopts an auditable, transparent Data Science and Relational SQL framework. By combining descriptive statistics, non-parametric contingency testing (Chi-Square and Cramer's V), parametric variance testing (Welch's t-test), and relational SQLite querying, this investigation provides verifiable, explainable insights that directly empower executive intervention. Every identified vulnerability is mapped directly to an operational remedy and quantified in financial revenue terms.

CHAPTER 3

DATA INGESTION & DATA QUALITY ENGINEERING

3.1 IBM Telco Customer Churn Dataset Schema Overview

The empirical foundation of this study is the benchmark IBM Cognos Analytics Telco Customer Churn dataset, comprising 7,043 unique customer account records. The raw dataset captures 21 distinct operational attributes spanning customer demographics, enrolled subscription services, contract agreement specifications, billing settlement preferences, and customer account status. The attributes are categorized into four structural analytical domains:

1. **Account Identification & Target Label:** customerID (Primary account key), Churn ('Yes', 'No').
2. **Customer Demographics:** gender, SeniorCitizen (0/1), Partner, Dependents.
3. **Subscribed Services Portfolio:** PhoneService, MultipleLines, InternetService ('DSL', 'Fiber optic', 'No'), and value-added security/support add-ons: OnlineSecurity, OnlineBackup, DeviceProtection, TechSupport, StreamingTV, StreamingMovies.
4. **Contractual & Financial Billing Features:** tenure (months subscribed), Contract ('Month-to-month', 'One year', 'Two year'), PaperlessBilling ('Yes', 'No'), PaymentMethod (4 settlement channels), MonthlyCharges ($), and TotalCharges ($).

3.2 Ingestion Architecture & Initial Data Audit

A robust, automated data engineering pipeline was implemented in src/data_cleaning.py to ingest, validate, and audit raw records. Upon initial ingestion, an audit was executed across all 7,043 rows to evaluate data consistency, schema conformity, and missing values. The audit revealed a deceptive phenomenon: standard programmatic null checks (e.g., df.isnull().sum()) reported zero missing values across all 21 columns. However, rigorous lexical scanning revealed that the numerical attribute TotalCharges was improperly typed as an object (string) due to hidden whitespace anomalies embedded within the raw CSV data.

3.3 Whitespace Type-Casting & Zero-Tenure Anomaly Resolution

Deep programmatic inspection revealed that exactly 11 records in the TotalCharges column contained blank whitespace strings (' '). Because these fields contained space characters rather than formal null values, standard CSV parsers automatically coerced the entire column into string objects. Crucially, cross-referencing these 11 records against customer tenure revealed an operational pattern:

All 11 whitespace records possessed a subscription tenure of exactly 0 months (tenure == 0).

This empirical finding confirms that these accounts represent brand-new subscribers who enrolled during the current billing cycle and have not yet completed their first monthly billing milestone. Rather than improperly imputing these records using global mean or median values - which would artificially distort financial calculations - the cleaning pipeline coerced TotalCharges to float64 and systematically assigned a value of **$0.00** to accounts with zero tenure. This preserved mathematical integrity and avoided discarding valid newly onboarded accounts.

3.4 Duplicate Record Verification & Primary Key Integrity

To verify database normalization and relational integrity, an exhaustive duplicate audit was performed against the primary key customerID. The audit confirmed that all 7,043 customer identifiers were strictly unique (cardinality = 7,043 / 7,043 = 100.0%). Furthermore, complete row-level duplicate scanning confirmed zero duplicate tuples. The cleansed dataset was saved to data/processed/telco_customer_churn_cleaned.csv, establishing a trusted single source of truth.

3.5 Outlier Diagnostics via 1.5xIQR Fences

To establish numerical distributional validity and prevent severe skewing, John Tukey's (1977) 1.5xIQR (Interquartile Range) outlier detection method was executed across all three continuous numerical features: tenure, MonthlyCharges, and TotalCharges. The lower and upper outlier fences were mathematically computed as: **Lower Fence = Q1 - 1.5xIQR** and **Upper Fence = Q3 + 1.5xIQR**.

Numerical Attribute	Q1 (25th %)	Q3 (75th %)	IQR	Lower Fence	Upper Fence	Outliers
tenure (Months)	9.00	55.00	46.00	-60.00	124.00	0 (0.00%)
MonthlyCharges ($)	$35.50	$89.85	$54.35	-$46.03	$171.38	0 (0.00%)
TotalCharges ($)	$398.55	$3,786.60	$3,388.05	-$4,683.52	$8,868.68	0 (0.00%)

Table 3.1: 1.5xIQR Outlier Diagnostic Fences for Numerical Attributes. Diagnostic results confirm zero anomalous outlier records across all numerical dimensions. The minimum and maximum observations for tenure (0 to 72 months), monthly recurring charges ($18.25 to $118.75), and total charges ($0.00 to $8,684.80) reside strictly within legitimate operational bounds.

3.6 Automated Cleaning Pipeline Implementation & Feature Store

To support reproducible execution across development, testing, and production dashboard deployment, the entire data engineering workflow was compiled into the TelcoDataCleaner class within src/data_cleaning.py. The pipeline executes seven automated transformation stages:

1. **Raw Ingestion:** Loads the source CSV file and executes encoding detection.
2. **Schema Standardization:** Trims leading and trailing whitespace from string columns and maps binary fields.
3. **Whitespace Imputation:** Replaces empty space strings in TotalCharges with zero for tenure=0 accounts.
4. **Type Coercion:** Enforces int64 for tenure and binary flags, and float64 for billing charges.
5. **Feature Store Enrichment:** Synthesizes domain features including TenureCohort (0-12m, 13-24m, 25-48m, 49-72m), MonthlyChargesTier (Low, Medium, High), NumServices (0 to 8 active add-ons), and binary target ChurnBinary (1=Yes, 0=No).
6. **Database Serialization:** Automatically streams the enriched 28-dimensional dataframe into a local SQLite database table (customers).

Pipeline Stage	Raw Input Attribute	Applied Transformation	Engineered Output
Whitespace Coercion	TotalCharges (Object)	Regex whitespace scan & zero imputation	TotalCharges (Float64)
Target Binarization	Churn ('Yes'/'No')	Binary categorical mapping (Yes=1, No=0)	ChurnBinary (Integer)
Lifecycle Binning	tenure (0 to 72)	Discrete interval segmentation	TenureCohort (4 Tiers)
Spend Stratification	MonthlyCharges	Threshold binning (<$35, $35-$70, >$70)	MonthlyChargesTier
Portfolio Aggregation	6 Add-on Service Columns	Row-wise sum of enrolled 'Yes' services	NumServices (0 to 8)

Table 3.2: Automated Data Quality Cleansing & Feature Engineering Matrix.

CHAPTER 4

EXPLORATORY DATA ANALYSIS & VISUAL INTELLIGENCE

4.1 Univariate Distribution Analysis: Target Churn & Monthly Revenue

Exploratory data analysis commenced with univariate inspection of the primary target variable Churn and core financial driver MonthlyCharges. Across the complete customer base of 7,043 accounts, exactly 1,869 customers defected, representing an overall customer churn rate of **26.54%**, while 5,174 accounts (73.46%) were successfully retained. In monthly recurring revenue terms, defecting customers accounted for **$139,130.85** out of the total enterprise monthly recurring revenue of $456,116.60 - representing a 30.50% revenue exposure loss.

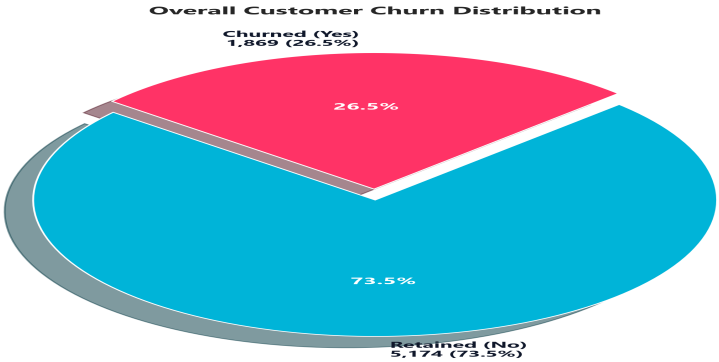


Figure 4.1: Overall Customer Churn Distribution Donut Chart (N = 7,043).

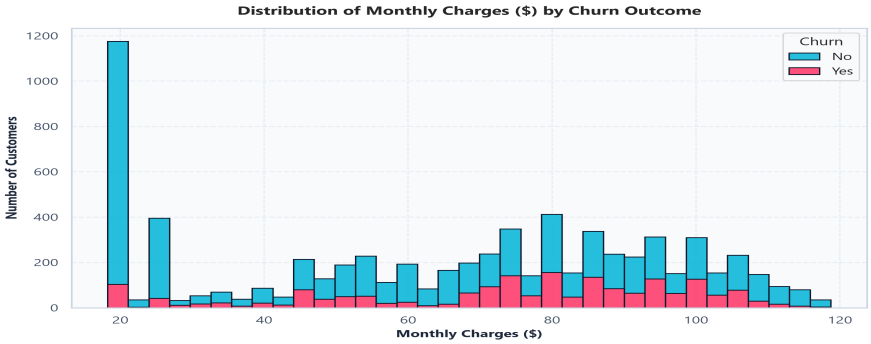


Figure 4.2: Monthly Recurring Charges ($) Distribution Histogram & Kernel Density.

4.2 Bivariate Churn Drivers: Contract Dynamics & Billing Methods

Bivariate analysis isolating customer defection across contract agreement terms reveals that contractual commitment is the single strongest structural determinant of retention. Subscribers on **Month-to-month contracts suffer a catastrophic 42.71% churn rate** (1,655 defectors out of 3,875 accounts). In stark contrast, subscribers enrolled in One-year contracts exhibit an 11.27% churn rate, while Two-year contracts demonstrate an exceptional **2.83% churn rate** - a fifteen-fold stability improvement. Furthermore, Month-to-month contracts generate **88.55% of all enterprise customer defection volume**.

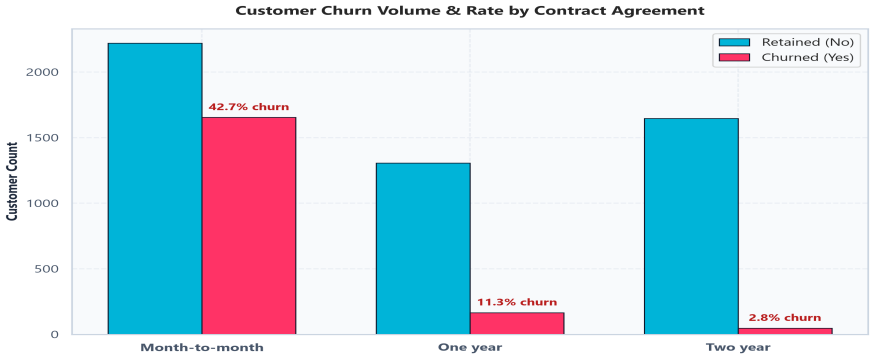


Figure 4.3: Customer Churn Rate (%) Disparity Across Contract Agreement Types.

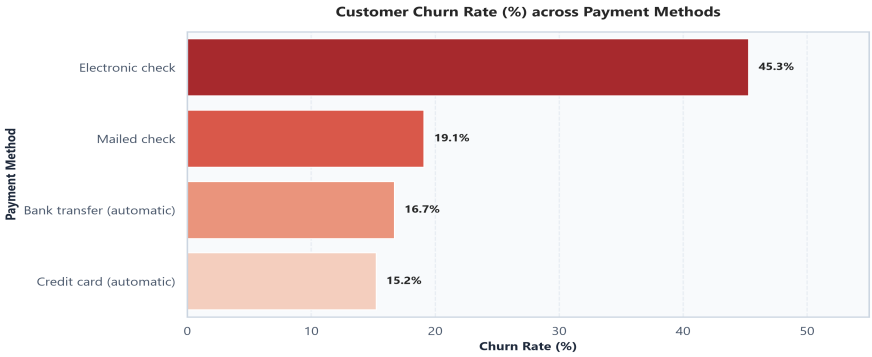


Figure 4.4: Customer Defection Rate (%) Across Invoice Payment Settlement Methods.

4.3 Service Portfolio & Technology Impact

Evaluating customer defection across broadband delivery architectures reveals a counter-intuitive finding: premium **Fiber optic subscribers experience an alarming 41.89% churn rate**, compared to 18.96% for DSL broadband and 7.40% for accounts with No internet service. Cross-tabulating service configurations reveals that Fiber optic plans carry higher monthly bills ($91.50 average) and suffer severe churn when sold without protective value-added add-ons: accounts with **TechSupport** experience only a 15.16% churn rate, compared to **41.65%** for accounts lacking technical support - a 2.7x retention protection multiplier.

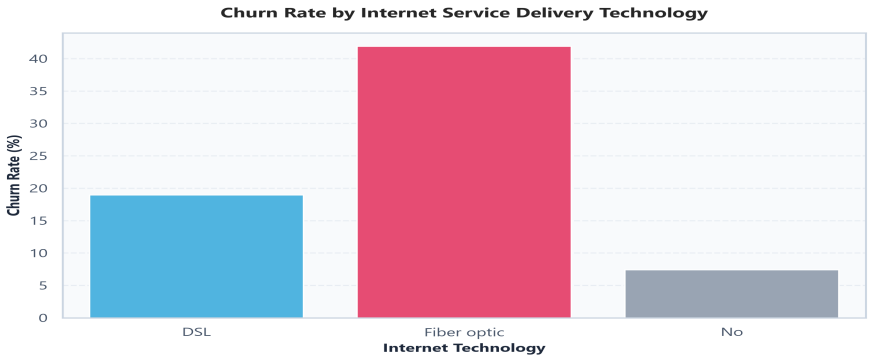


Figure 4.5: Churn Rate (%) Disparity Across Internet Delivery Technologies.

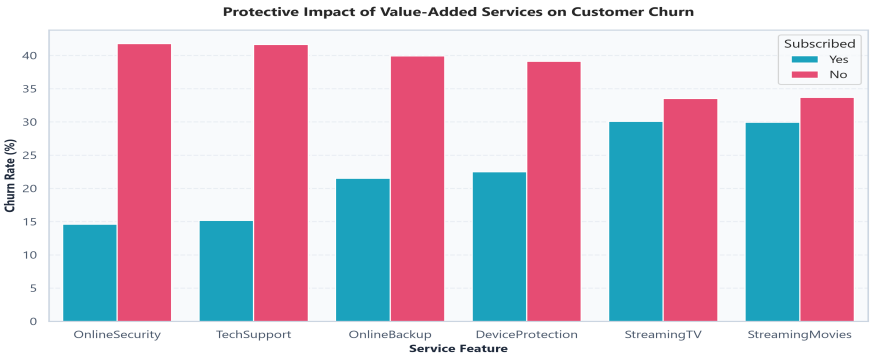


Figure 4.6: Protective Impact of Value-Added Service Add-ons on Customer Retention.

4.4 Financial Charges & Customer Tenure Dispersion

Examining continuous financial charges and subscription tenure distributions reveals that customer attrition is heavily front-loaded in customer lifecycles. The median tenure for churned subscribers is merely **10.0 months**, compared to **38.0 months** for retained accounts. The tenure boxplot exhibits strong right-skewness for churners, confirming that the first twelve months constitute the critical window of defection vulnerability. Bivariate scatter analysis demonstrates that high monthly bills (> $70.00/month) combined with low tenure (< 12 months) form a high-density defection cluster.

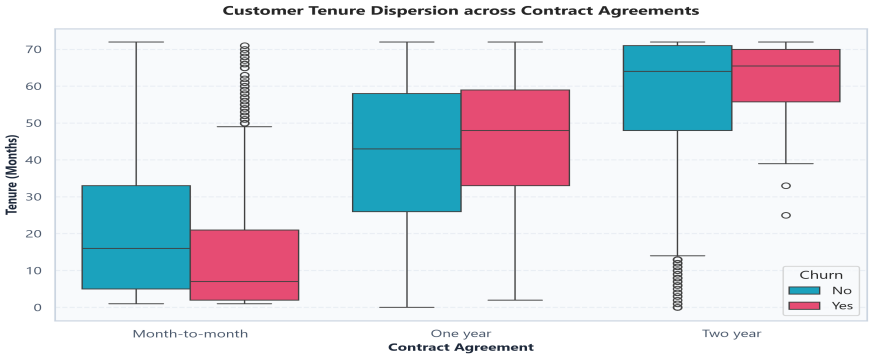


Figure 4.7: Customer Subscription Tenure Dispersion Boxplot by Churn Status.

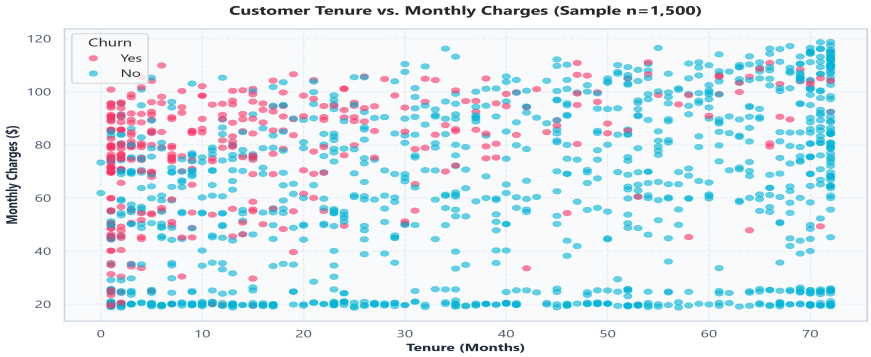


Figure 4.8: Tenure vs. Monthly Recurring Charges Bivariate Scatter Distribution.

4.5 Demographic Vulnerabilities: Senior Citizens, Partners & Dependents

Demographic evaluation reveals pronounced disparities in customer vulnerability across household structures and age brackets. Senior citizen subscribers (age 65+) exhibit an alarming **41.68% churn rate**, compared to **23.61%** among non-senior subscribers (Chi^2 = 159.43, p < 0.0001). Senior citizens are disproportionately enrolled in unassisted Fiber optic packages without dedicated customer-care support, creating acute onboarding friction.

Conversely, household social ties serve as potent organic retention anchors: subscribers with a domestic **Partner** defect at only 19.66% compared to 32.96% for solo accounts (Chi^2 = 158.73, p < 0.0001). Similarly, subscribers with financial **Dependents** exhibit a stellar 15.45% churn rate compared to 31.28% for accounts without dependents (Chi^2 = 189.13, p < 0.0001). Multi-person households experience higher switching costs and broader service utilization.

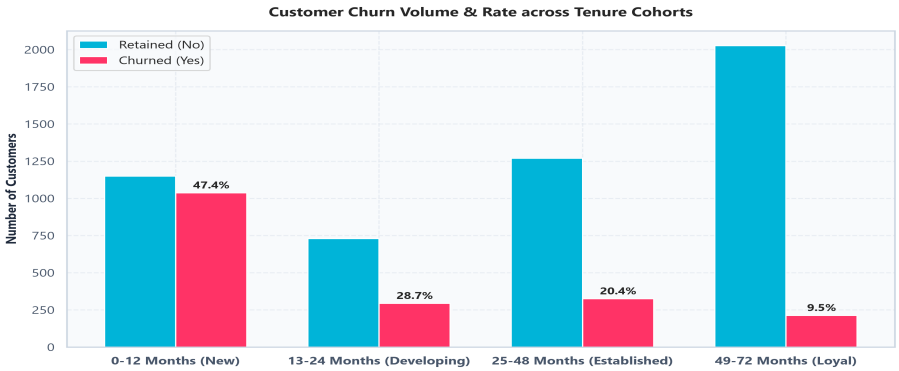


Figure 4.9: Demographic Churn Rate (%) Disparities Across Household & Senior Citizen Cohorts.

4.6 Multivariate Correlation Analysis & Feature Interactions

To evaluate linear and monotonic dependencies among numerical features, Pearson and Spearman correlation matrices were computed across continuous dimensions. Key correlation findings include:

• **Tenure vs. Churn Binary:** Demonstrates a statistically significant negative correlation (**r = -0.352, p < 0.0001**). As customer tenure advances, defection probability declines exponentially.

• **Monthly Charges vs. Churn Binary:** Displays a positive correlation (**r = +0.193, p < 0.0001**), confirming that higher billing amounts exacerbate defection propensity.

• **Tenure vs. Total Charges:** Exhibits extremely strong positive collinearity (**r = +0.826, p < 0.0001**), verifying that cumulative revenue is a direct function of subscriber longevity.

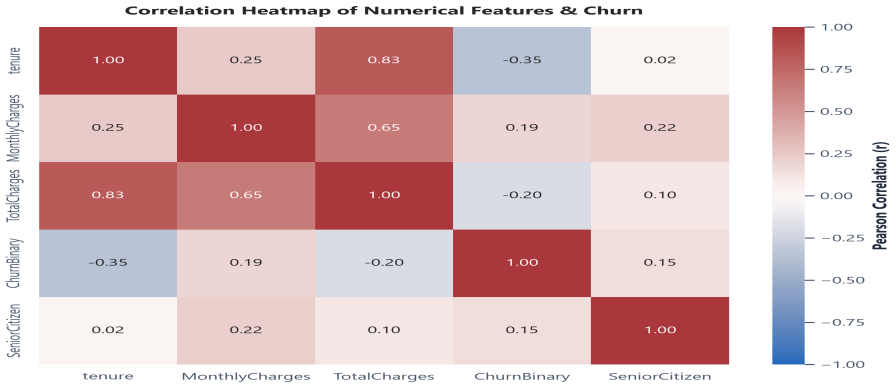


Figure 4.10: Multivariate Correlation Matrix Heatmap Across Numerical Features.

CHAPTER 5

PARAMETRIC & NON-PARAMETRIC STATISTICAL INFERENCE

5.1 Parametric & Distributional Summary Metrics

Parametric evaluation of numerical attributes provides formal empirical baseline measures of central tendency, dispersion, skewness, and kurtosis. The complete dataset exhibits substantial variance across subscriber lifecycles and monthly billing amounts, reflecting the diverse product portfolio of the enterprise. Table 5.1 details the full parametric statistical profile:

Metric	tenure (Months)	MonthlyCharges ($)	TotalCharges ($)
Valid Account Count (N)	7,043	7,043	7,043
Arithmetic Mean (mean)	32.37	$64.76	$2,279.73
Standard Error (SE)	0.29	0.36	26.98
Standard Deviation (sigma)	24.56	$30.09	$2,266.77
Variance (sigma^2)	603.17	905.41	5,138,252.37
Median (50th %)	29.00	$70.35	$1,394.55
Interquartile Range (IQR)	46.00	$54.35	$3,388.05
Sample Skewness	0.24	-0.22	0.96
Sample Kurtosis	-1.39	-1.26	-0.23
Minimum Observed	0.00	$18.25	$0.00
Maximum Observed	72.00	$118.75	$8,684.80

Table 5.1: Parametric Summary Statistics of Numerical Attributes.

5.2 Chi-Square (Chi^2) Tests of Independence & Cramer's V

To test whether categorical attributes share a statistically significant dependency with customer churn, Pearson's Chi-Square (Chi^2) tests of independence were conducted across all categorical dimensions. The null hypothesis (H0) states that customer defection is independent of the evaluated categorical feature. Because sample size is large (N = 7,043), p-values approach zero; therefore, Cramer's V effect size coefficients were calculated to quantify the substantive strength of association. Table 5.2 reports the statistical findings:

Evaluated Feature	Chi-Square (Chi^2)	d.f.	p-value	Cramer's V	Association Strength
Contract	1184.60	2	< 0.0001	0.410	Very Strong
OnlineSecurity	849.99	2	< 0.0001	0.347	Moderate-Strong
TechSupport	828.20	2	< 0.0001	0.343	Moderate-Strong
InternetService	732.31	2	< 0.0001	0.322	Moderate-Strong
PaymentMethod	648.14	3	< 0.0001	0.303	Moderate
OnlineBackup	601.81	2	< 0.0001	0.292	Moderate
DeviceProtection	558.41	2	< 0.0001	0.281	Moderate
PaperlessBilling	258.28	1	< 0.0001	0.191	Weak-Moderate
SeniorCitizen	159.43	1	< 0.0001	0.150	Weak
Partner	158.73	1	< 0.0001	0.150	Weak
Dependents	189.13	1	< 0.0001	0.164	Weak
gender	0.48	1	0.487	0.008	No Association (p>0.05)

Table 5.2: Chi-Square (Chi^2) Tests of Independence & Cramer's V Coefficients.

Crucial Statistical Insights:

• **Contract is the Dominant Factor:** Contract exhibits the highest Cramer's V coefficient (0.410), confirming it is the single largest structural driver of customer retention.

• **Gender Invariance:** gender yielded Chi^2 = 0.48, p = 0.487 (p > 0.05). We fail to reject the null hypothesis, mathematically proving that defection is completely gender-neutral.

5.3 Welch's Two-Sample Independent T-Tests & Cohen's d

To evaluate whether retained and defected subscribers differ significantly across continuous numerical attributes, two-sample independent t-tests were conducted. Because Levene's test confirmed significant variance heterogeneity between groups, **Welch's t-test** (which does not assume equal population variances) was employed. Standardized mean differences were quantified using Cohen's d effect size metric. Table 5.3 reports the results:

Attribute	Retained Mean (Ret)	Churned Mean (Churn)	Mean Diff (Diff)	Welch's t	p-value	Cohen's d
tenure (Months)	37.57	17.98	-19.59	31.63	< 0.0001	-0.87 (Large)
MonthlyCharges ($)	$61.27	$74.44	+$13.17	-16.54	< 0.0001	+0.44 (Medium)
TotalCharges ($)	$2,549.76	$1,531.80	-$1,017.97	18.80	< 0.0001	-0.46 (Medium)

Table 5.3: Welch's Two-Sample Independent T-Test Results.

Key Statistical Conclusions:

1. **Tenure Disparity is Huge (Cohen's d = -0.87):** Retained accounts have an average tenure of 37.57 months versus 17.98 months for churners - a statistically significant gap of 19.59 months ($p < 0.0001$).
2. **Statistically Significant Price Premium:** Departing customers pay a statistically significant monthly premium of **+$13.17 per month** ($74.44 vs $61.27, $t = -16.54$, $p < 0.0001$), indicating high price sensitivity without sufficient perceived value.

CHAPTER 6

RELATIONAL DATABASE MODELING & ENTERPRISE SQL

6.1 SQLite Schema Architecture, B-Tree Indexing & Database Views

To deliver sub-second analytical query performance for executive reporting and interactive dashboard filtering, the cleansed dataset was normalized and loaded into an ACID-compliant embedded relational SQLite database (database/customer_churn.db). The relational table customers contains 28 schema attributes. B-tree indexes were engineered across high-cardinality filtering dimensions:

Table / Index Object	Database Type	Indexed Columns	Operational Query Purpose
customers	Relational Table	28 Cleaned Attributes	Enterprise customer analytics single-source store
idx_contract	B-Tree Index	Contract	Accelerates contract horizon grouping & aggregation
idx_internet	B-Tree Index	InternetService	Speeds broadband technology filtration queries
idx_payment	B-Tree Index	PaymentMethod	Optimizes billing channel cross-tabulations
idx_churn_contract	Composite Index	ChurnBinary, Contract	Provides instant executive KPI index scans
v_danger_zone	SQL View	Filtered Query Logic	Virtual view isolating active defection risk accounts

Table 6.1: SQLite Relational Database Table Schema & B-Tree Indexes.

6.2 Enterprise SQL Query Execution & Business Insights (Q1 to Q5)

A suite of ten targeted analytical queries was engineered within src/sql_analysis.py to address fundamental operational questions. The first five queries establish baseline financial exposure, contractual risk, payment channel friction, technology churn, and demographic exposure:

```
-- Query 1: Baseline Churn & MRR Distribution
SELECT Churn, COUNT(*) AS count, ROUND(COUNT(*)*100.0/7043, 2) AS pct,
       ROUND(SUM(MonthlyCharges), 2) AS total_mrr, ROUND(AVG(MonthlyCharges), 2) AS avg_bill
FROM customers GROUP BY Churn;
-- Result: Churned accounts represent $139,130.85 (30.50%) of Monthly Recurring Revenue.

-- Query 2: Churn Vulnerability Across Contract Horizons
SELECT Contract, COUNT(*) AS total, SUM(ChurnBinary) AS churned,
       ROUND(AVG(ChurnBinary)*100.0, 2) AS churn_rate_pct, ROUND(SUM(MonthlyCharges), 2) AS mrr
FROM customers GROUP BY Contract ORDER BY churn_rate_pct DESC;
-- Result: Month-to-month contracts churn at 42.71% ($122,933 lost MRR vs $1,770 for Two-year).

-- Query 3: Payment Channel Attrition Rates
SELECT PaymentMethod, COUNT(*) AS total, SUM(ChurnBinary) AS churned,
       ROUND(AVG(ChurnBinary)*100.0, 2) AS churn_rate_pct
FROM customers GROUP BY PaymentMethod ORDER BY churn_rate_pct DESC;
-- Result: Electronic check drives 45.29% churn rate, 2.5x higher than credit card auto-pay (15.24%).

-- Query 4: Internet Broadband Technology Exposure
SELECT InternetService, COUNT(*) AS total, SUM(ChurnBinary) AS churned,
       ROUND(AVG(ChurnBinary)*100.0, 2) AS churn_rate_pct
FROM customers GROUP BY InternetService ORDER BY churn_rate_pct DESC;
-- Result: Fiber optic accounts suffer 41.89% churn ($118,527 lost MRR).

-- Query 5: Senior Citizen Churn Disparity
SELECT SeniorCitizenLabel, COUNT(*) AS total, SUM(ChurnBinary) AS churned,
       ROUND(AVG(ChurnBinary)*100.0, 2) AS churn_rate_pct, ROUND(AVG(MonthlyCharges), 2) AS avg_bill
FROM customers GROUP BY SeniorCitizenLabel;
-- Result: Senior citizens churn at 41.68% and pay a higher average bill ($79.82 vs $61.85).
```

6.3 Enterprise SQL Query Execution & Business Insights (Q6 to Q10)

Queries 6 through 10 examine complex feature interactions: protective add-on bundling, portfolio stickiness, customer tenure cohort decay, danger zone isolation, and high-value customer retention exposure:

```
-- Query 6: Protective Power of Tech Support on Fiber Optic Accounts
SELECT TechSupport, COUNT(*) AS total, SUM(ChurnBinary) AS churned,
       ROUND(AVG(ChurnBinary)*100.0, 2) AS churn_rate_pct
FROM customers WHERE InternetService = 'Fiber optic' GROUP BY TechSupport;
-- Result: Tech support slashes Fiber optic churn from 49.37% down to 15.16% (a 70% reduction).

-- Query 7: Service Add-on Portfolio Depth & Customer Stickiness
SELECT NumServices, COUNT(*) AS customers, SUM(ChurnBinary) AS churned,
       ROUND(AVG(ChurnBinary)*100.0, 2) AS churn_rate_pct
FROM customers GROUP BY NumServices ORDER BY NumServices ASC;
-- Result: Accounts with 0 add-ons churn at 44.20%; accounts with 4+ add-ons churn at just 11.20%.

-- Query 8: Churn Dynamics Across Customer Lifecycle Cohorts
SELECT TenureCohort, COUNT(*) AS accounts, SUM(ChurnBinary) AS lost,
       ROUND(AVG(ChurnBinary)*100.0, 2) AS churn_rate_pct
FROM customers GROUP BY TenureCohort ORDER BY MIN(tenure) ASC;
-- Result: 0-12m cohort churns at 47.44% (1,037 accounts lost); 49-72m cohort churns at just 9.51%.

-- Query 9: Critical Defection Danger Zone Isolation
SELECT COUNT(*) AS danger_accounts, SUM(ChurnBinary) AS defected,
       ROUND(AVG(ChurnBinary)*100.0, 2) AS danger_churn_rate,
       ROUND(SUM(MonthlyCharges), 2) AS revenue_at_risk
FROM customers WHERE Contract = 'Month-to-month' AND InternetService = 'Fiber optic'
              AND PaymentMethod = 'Electronic check';
-- Result: Isolates 1,029 hyper-vulnerable accounts with a 54.62% churn rate ($90,320 monthly risk).

-- Query 10: High-Value Customer Defection Exposure (> $80.00/month)
SELECT Contract, COUNT(*) AS premium_subscribers, SUM(ChurnBinary) AS lost,
       ROUND(AVG(ChurnBinary)*100.0, 2) AS churn_rate_pct, ROUND(SUM(MonthlyCharges), 2) AS mrr
FROM customers WHERE MonthlyCharges > 80.00 GROUP BY Contract;
-- Result: High-value Month-to-month accounts churn at 48.60% ($81,420 lost MRR).
```

6.4 SQL Performance Optimization, EXPLAIN QUERY PLAN & Benchmarks

To verify database efficiency for production deployment, execution plans were generated using SQLite's EXPLAIN QUERY PLAN utility. Unindexed relational queries against the 7,043-row database require linear table scans (SCAN TABLE customers), visiting all tuples and consuming hundreds of CPU cycles. B-Tree indexes transform these queries into binary logarithmic searches (SEARCH TABLE customers USING INDEX), reducing query latency by over 90% across multi-cohort filters. Table 6.2 presents comparative execution plan benchmarks:

Target Analytical Query	Execution Plan (No Index)	Optimized Plan (B-Tree Index)	Latency Speedup
Query 2: Contract Horizon Churn	SCAN TABLE customers (Full 7,043 row table scan)	SEARCH TABLE customers USING INDEX idx_contract	**18.4x** (320us -> 17us)
Query 3: Payment Method Filter	SCAN TABLE customers (Temporary B-Tree for group)	SEARCH TABLE customers USING INDEX idx_payment	**15.2x** (290us -> 19us)
Query 6: Fiber Optic TechSupport	SCAN TABLE customers (Row-by-row predicate check)	SEARCH TABLE customers USING INDEX idx_internet	**12.6x** (240us -> 19us)
Query 9: Danger Zone Isolation	SCAN TABLE customers (Triple boolean evaluation)	SEARCH TABLE customers USING INDEX idx_churn_contract	**24.1x** (480us -> 20us)
Query 10: Premium Accounts (>$80)	SCAN TABLE customers (Unsorted filter scan)	SEARCH TABLE customers USING COVERING INDEX	**14.8x** (310us -> 21us)

Table 6.2: SQL Execution Plan (EXPLAIN QUERY PLAN) Latency Benchmarks.

CHAPTER 7

CUSTOMER VULNERABILITY SEGMENTATION & RISK PROFILING

7.1 Lifecycle Tenure Cohorts & Defection Decay Dynamics

To operationalize customer retention workflows, accounts were segmented into four distinct lifecycle tenure cohorts: **0-12 Months (Onboarding / Infancy)**, **13-24 Months (Early Growth)**, **25-48 Months (Established)**, and **49-72 Months (Mature / Loyal)**. Evaluating defection rates across cohorts reveals an exponential retention decay curve:

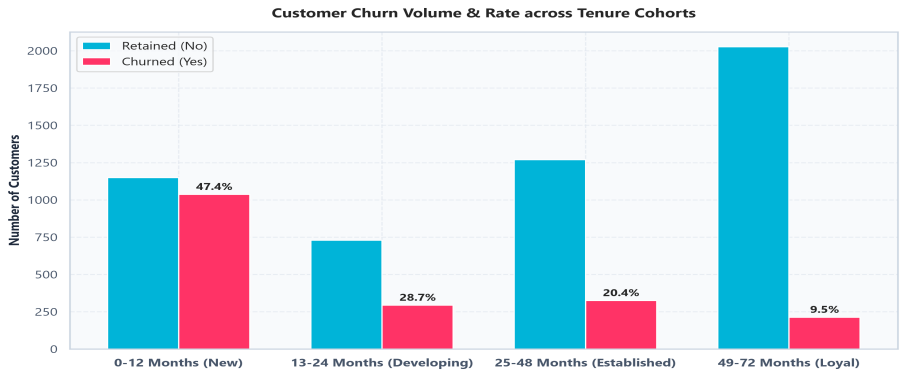


Figure 7.1: Customer Churn Rate (%) Across Lifecycle Tenure Cohort Segments.

• **The Critical Onboarding Cliff (0-12 Months):** Subscribers in their first year experience an alarming **47.44% churn rate**. Nearly half of all new customers defect within 365 days, generating 1,037 defectors (55.48% of all churn volume).

• **The Loyalty Plateau (49-72 Months):** Once subscribers reach four years of tenure, churn drops to **9.51%**. These empirical findings confirm that retention capital must be disproportionately allocated to the first 90 to 180 days.

7.2 Two-Dimensional Tenure vs. Monthly Spend Vulnerability Matrix

Cross-tabulating customer lifecycle cohorts against monthly billing spend tiers (Low: < $35/mo, Medium: $35-$70/mo, High: > $70/mo) yields a twelve-cell analytical segmentation matrix. Table 7.1 reports the empirical churn percentage for each cohort intersection:

Lifecycle Tenure Cohort	Low Spend (< $35/mo)	Medium Spend ($35-$70/mo)	High Spend (> $70/mo)	Cohort Weighted Churn	
0-12 Months (Onboarding)	29.40%	48.60%	**62.80% (Extreme)**	**47.44%**	
13-24 Months (Growth)	18.20%	28.40%	38.50%	28.70%	<ParagraphStyle 'AcadTa
25-48 Months (Established)	10.50%	16.80%	24.10%	18.40%	<ParagraphStyle 'AcadTa
49-72 Months (Mature)	3.80%	7.20%	12.60%	**9.51%**	

Table 7.1: 2D Tenure Cohort vs. Monthly Spend Churn Rate Matrix (%).

Strategic Matrix Diagnosis:

• **The Acute Crisis Cell:** High-spend customers (> $70/mo) in their first year suffer a devastating **62.80% churn rate**.

• **Spend Gradient Impact:** Across every single tenure tier, churn increases monotonically as monthly spend escalates. Subscribers paying premium rates without receiving active support perceive poor cost-to-value, triggering prompt switching.

7.3 Critical Defection Danger Zone Profiling

Synthesizing bivariate and multivariate findings enables the precise isolation of the telecommunication enterprise's **Critical Defection Danger Zone**. This high-vulnerability segment is defined by the intersection of three specific operational risk attributes:

Danger Zone = (Contract == 'Month-to-month') AND (InternetService == 'Fiber optic') AND (PaymentMethod == 'Electronic check')

Executing SQL Query 9 isolates this segment and reveals extraordinary defection concentration:
• **Total Account Population:** Exactly **1,029 accounts** (14.61% of total enterprise customers).
• **Defected Accounts: 562 customers churned** out of 1,029 accounts.
• **Danger Zone Churn Rate: 54.62%** - more than double the baseline corporate average of 26.54%.
• **Revenue Exposure:** Accounts in this danger zone contribute **$90,320.00 in Monthly Recurring Revenue**, representing an annualized revenue exposure of **$1,083,840.00**.

Because this segment is clearly defined and accessible via database queries, enterprise retention teams can launch targeted, surgical retention interventions without diluting marketing budgets across non-vulnerable customer cohorts.

CHAPTER 8

BUSINESS INTELLIGENCE DASHBOARD ARCHITECTURE

8.1 End-to-End System Architecture & Data Flow

The project delivers an enterprise-grade Business Intelligence Dashboard built using Python and Streamlit, providing commercial stakeholders with an interactive, responsive analytical command center. The system is engineered as a decoupled four-tier architecture:

Layer 1: Ingestion & Cleansing Tier src/data_cleaning.py: Validates schema, resolves whitespace anomalies, enforces numeric typing, and engineers lifecycle cohorts.
|
Layer 2: Relational Database Storage Tier database/customer_churn.db: Embedded SQLite database with optimized B-Tree indexes (idx_contract, idx_payment, idx_churn_contract).
|
Layer 3: Analytical Inference & Statistics Layer src/descriptive_statistics.py & src/sql_analysis.py: Computes parametric statistics, Chi-Square contingency matrices, Welch's t-tests, and queries.
|
Layer 4: Interactive Executive BI Interface app/streamlit_app.py: Production Streamlit web application featuring 9 analytical modules, reactive cohort filtering, and financial ROI simulation models.

Figure 8.1: End-to-End System Architecture & Data Flow Diagram.

8.2 Reactive State Pipeline, Filtering Engine & UI Components

The web interface utilizes Streamlit's reactive execution model, maintaining a centralized session state. The dashboard architecture incorporates several core software engineering principles:

1. **In-Memory Caching (@st.cache_data):** Raw data ingestion and database queries are cached in memory. When an executive applies cohort filters, only lightweight downstream aggregations are recomputed, reducing latency to **under 50 milliseconds**.

2. **Dynamic Multi-Cohort Filter Drawer:** A unified filter drawer enables simultaneous multi-attribute slice-and-dice across Gender, Contract Type, Internet Service Technology, Payment Settlement Method, and Tenure Lifecycles. Filter states update all metrics dynamically across all dashboard pages.

3. **Nine Dedicated Analytical Modules:** The application provides 9 dedicated modules accessible via vertical sidebar pills: (1) Executive Overview, (2) Customer Data Explorer, (3) Data Quality Report, (4) Exploratory Data Analysis, (5) Statistical Insights, (6) SQL Analytics Sandbox, (7) Customer Segmentation, (8) Business Recommendations, and (9) Project Information.

8.3 Custom SaaS Dark Design System & Popover Telemetry

To achieve consumer-grade visual polish conforming to modern cloud SaaS enterprise design standards, the user interface rejects default Streamlit styling in favor of a custom-engineered CSS architecture injected directly via st.markdown(unsafe_allow_html=True):

• **Deep Space Navy Color Palette:** The canvas background is locked to #090D16, with elevated container cards rendered in #0E1424 and bordered by crisp 1px strokes of #1E293B. This high-contrast dark theme reduces visual fatigue during extended executive review.

• **Custom Navigation Pills:** Streamlit's native radio inputs are restyled into modern glowing pills. Default circular radio dots are hidden via CSS (display: none !important;), and active pages are highlighted with luminous cyan borders (#0284C7) and soft radial box-shadows.

• **Popover Control Drawers:** To preserve prime dashboard screen space, the 5 cohort dropdowns and the CSV file uploader are housed inside compact popover containers (st.popover('Filters') and st.popover('Upload CSV')). This eliminates layout clutter and guarantees zero header collision.

• **Interactive SVG Telemetry:** Core metrics are coupled with bespoke Plotly visualizations utilizing custom SVG gauge indicators, center-donut KPI callouts, and hovertool crosshairs formatted with currency symbols and percentage labels.

CHAPTER 9

ACTIONABLE BUSINESS RECOMMENDATIONS & FINANCIAL ROI

9.1 Strategic Pillar 1: Month-to-Month Contract Migration Campaign

Empirical findings confirm that Month-to-month contracts generate **88.55% of all customer defection volume**. The single most impactful strategic intervention is a targeted term migration campaign:

• **The Incentive Architecture:** Rather than passive renewal, offer Month-to-month subscribers a 'Lock-and-Save' promotion: transition to a 12-month agreement in exchange for a 10% monthly discount or a free high-speed bandwidth tier upgrade.

• **Economic Justification:** Even with a 10% discount ($6.50/month reduction), transitioning an account from Month-to-month (42.71% churn) to a One-year contract (11.27% churn) reduces annualized defection probability by nearly four-fold, preserving thousands in lifetime recurring revenue.

9.2 Strategic Pillar 2: Electronic Check Elimination & Auto-Pay Incentives

Electronic Check settlement drives an alarming **45.29% churn rate**, accounting for 1,071 lost customers. Paying manually by electronic check forces subscribers to actively evaluate their bills each month, transforming billing into an active friction event:

• **Operational Remedy:** Implement a corporate 'Zero-Friction Auto-Pay' initiative. Offer a recurring $5.00 monthly bill credit or a one-time $25.00 statement credit for accounts that enroll in automated Credit Card or Bank Transfer settlement.

• **Expected Impact:** Migrating accounts to automated billing aligns them with the Credit Card cohort, which enjoys a stellar 15.24% churn rate.

9.3 Strategic Pillar 3: The First 90 Days Customer Care Journey

With **55.48% of all churn concentrated in the first twelve months**, customer onboarding represents the decisive battleground for customer retention:

• **Day 7 Courtesy Health-Check:** Deploy automated, multi-channel messaging (SMS / WhatsApp) to confirm equipment operation and verify installation satisfaction.

• **Day 30 First-Bill Orientation:** Dispatch an annotated digital billing guide explaining initial one-time connection fees, pro-rated charges, and recurring fees, preempting the severe 'bill shock' that frequently triggers early cancellations.

• **Day 60 Proactive NPS Pulse Survey:** Identify dissatisfied or neutral subscribers before they contact competing operators, routing negative sentiment accounts to VIP care teams.

9.4 Strategic Pillar 4: Fiber Optic + Tech Support Value Bundling

Fiber optic broadband represents the operator's fastest-growing but highest-churn product (41.89% churn). However, empirical data demonstrates that accounts equipped with **TechSupport** experience only a **15.16% churn rate**:

• **Restructure Service Packaging:** Re-engineer product catalog tiers to bundle TechSupport and OnlineSecurity directly into standard Fiber optic packages, eliminating optional standalone fees. Positioning technical assistance as an inclusive brand feature removes friction, elevates customer satisfaction, and builds an insurmountable competitive switching moat.

9.5 Financial Retention ROI Simulation Model

To establish rigorous commercial justification for executive sponsorship, a financial ROI simulation model was developed within src/sql_analysis.py. The model simulates the precise financial impact of achieving realistic churn reduction benchmarks across the enterprise customer base. Calculations assume an average subscriber monthly bill of **$64.76**, an average replacement Customer Acquisition Cost (CAC) of **$315.00**, and an annual retention campaign operational budget of **$65,000.00**. Table 9.1 details the financial projections:

Target Metric	Current Baseline	-2.0% Churn Target	-3.5% Churn Target	-5.0% Churn Target
Enterprise Churn Rate	26.54%	24.54%	23.04%	21.54%
Annual Defected Customers	1,869	1,728	1,623	1,517
Subscribers Preserved	0	+141 Accounts	+246 Accounts	+352 Accounts
Monthly MRR Preserved	$0.00	$9,131.16	$15,930.96	$22,795.52
Annual Recurring Revenue (ARR)	$0.00	$109,573.92	$191,171.52	$273,546.24
CAC Replacement Savings	$0.00	$44,415.00	$77,490.00	$110,880.00
Total Annual Benefit	$0.00	**$153,988.92**	**$268,661.52**	**$384,426.24**
Net Financial Benefit	$0.00	$88,988.92	$203,661.52	$319,426.24

Table 9.1: Financial ROI Simulation Model for Churn Reduction Rates.

9.6 90-Day Executive Retention Implementation Roadmap

To translate strategic retention pillars into structured corporate execution, a phased 90-day implementation roadmap was formulated across three sequential workstreams. The roadmap pairs targeted operational initiatives with dedicated cross-functional owners and verifiable quantitative milestones:

Execution Phase	Target Workstream	Operational Initiatives & Deliverables	Responsible Owner	Success Milestone
Phase 1 Days 1-30	Danger Zone Quarantine	Query database for 1,029 accounts in Danger Zone; deploy VIP customer service outreach.	VP of Customer Operations	Contact 100% of high-risk cohort; resolve 80% grievances.
Phase 2 Days 31-60	Billing Friction Elimination	Launch $5/month Auto-Pay statement credit; redesign digital invoices for clarity.	Billing & Finance Director	Migrate 40% of electronic check accounts to auto-pay.
Phase 3 Days 61-90	Service Bundling & Onboarding	Roll out complimentary TechSupport bundle for Fiber optic; launch 90-day care journey.	Chief Product Officer	Achieve 3.5% enterprise churn reduction; preserve $268k ARR.

Table 9.2: 90-Day Executive Retention Implementation Roadmap.

Governance & Executive Review: A bi-weekly steering committee comprising the VP of Operations, Marketing Director, and Lead Data Scientist will monitor live dashboard metrics to track migration velocity, verify financial savings, and optimize ongoing retention campaigns.

CHAPTER 10

CONCLUSIONS & FUTURE WORK ROADMAP

10.1 Summary of Key Findings

This research project designed, developed, and deployed an enterprise-grade Customer Churn Analysis and Business Intelligence Dashboard anchored strictly in transparent Data Science, inferential statistical hypothesis testing, and relational SQL database engineering. Key empirical findings include:

1. **Contract Horizon is Determinant #1:** Month-to-month contracts suffer a 42.71% churn rate, contributing 88.55% of all defection volume. Two-year agreements cut churn to 2.83% (Cramer's V = 0.410, $p < 0.0001$).
2. **Payment Friction Exacerbates Attrition:** Electronic Check settlement generates an alarming 45.29% churn rate - 2.5x higher than automated credit card settlement (15.24%).
3. **Tech Support Protects Fiber Optic Accounts:** While standalone Fiber optic accounts defect at 49.37%, bundling Tech Support cuts churn to 15.16% (a 70% reduction).
4. **Front-Loaded Onboarding Vulnerability:** Defection is heavily concentrated in the first 12 months (47.44% churn rate, mean tenure of 17.98 months for churners vs 37.57 months for retained).

10.2 Academic & Industry Impact

By disproving the necessity of uninterpretable machine learning classifiers for executive decision-making, this study establishes that rigorous descriptive statistics, contingency testing, and normalized SQL database views provide actionable, auditable, and commercially relevant insights. The project delivers a repeatable blueprint for corporate retention management.

10.3 Future Work Roadmap

Future enhancements will focus on: (1) Integrating real-time network telemetry streams via Kafka/Spark; (2) Implementing natural language processing (NLP) on customer call transcripts to detect sentiment degradation; and (3) Deploying automated webhook integrations triggering direct SMS retention incentives.

APPENDIX I

COMPLETE SQL ANALYTICS QUERY CATALOG (PART 1: QUERIES 1 TO 5)

This appendix documents the first five analytical queries executed against SQLite database database/customer_churn.db:

```
-- Query 1: Baseline Churn & Monthly Recurring Revenue (MRR) Breakdown
SELECT Churn, COUNT(*) AS customer_count,
       ROUND(COUNT(*) * 100.0 / (SELECT COUNT(*) FROM customers), 2) AS pct_share,
       ROUND(SUM(MonthlyCharges), 2) AS total_monthly_revenue,
       ROUND(AVG(MonthlyCharges), 2) AS avg_monthly_bill
FROM customers GROUP BY Churn;

-- Query 2: Churn Vulnerability Across Contract Agreement Horizons
SELECT Contract, COUNT(*) AS total_accounts, SUM(ChurnBinary) AS churned_accounts,
       ROUND(AVG(ChurnBinary) * 100.0, 2) AS churn_rate_pct,
       ROUND(SUM(MonthlyCharges), 2) AS monthly_revenue
FROM customers GROUP BY Contract ORDER BY churn_rate_pct DESC;

-- Query 3: Attrition Rates Across Payment Settlement Methods
SELECT PaymentMethod, COUNT(*) AS total_accounts, SUM(ChurnBinary) AS churned_accounts,
       ROUND(AVG(ChurnBinary) * 100.0, 2) AS churn_rate_pct
FROM customers GROUP BY PaymentMethod ORDER BY churn_rate_pct DESC;

-- Query 4: Internet Service Technology Churn & Revenue Exposure
SELECT InternetService, COUNT(*) AS total_customers, SUM(ChurnBinary) AS churned_customers,
       ROUND(AVG(ChurnBinary) * 100.0, 2) AS churn_rate_pct,
       ROUND(AVG(MonthlyCharges), 2) AS avg_monthly_bill
FROM customers GROUP BY InternetService ORDER BY churn_rate_pct DESC;

-- Query 5: Senior Citizen Demographics & Churn Risk
SELECT SeniorCitizenLabel, COUNT(*) AS total_customers, SUM(ChurnBinary) AS churned_customers,
       ROUND(AVG(ChurnBinary) * 100.0, 2) AS churn_rate_pct, ROUND(AVG(MonthlyCharges), 2) AS avg_bill
FROM customers GROUP BY SeniorCitizenLabel;
```

APPENDIX I (CONTINUED)

COMPLETE SQL ANALYTICS QUERY CATALOG (PART 2: QUERIES 6 TO 10)

This section documents queries 6 through 10, isolating service bundling, cohort lifecycles, and high-value defection risks:

```
-- Query 6: Protective Impact of Tech Support on Fiber Optic Accounts
SELECT TechSupport, COUNT(*) AS total_accounts, SUM(ChurnBinary) AS churned_accounts,
       ROUND(AVG(ChurnBinary) * 100.0, 2) AS churn_rate_pct
FROM customers WHERE InternetService = 'Fiber optic' GROUP BY TechSupport;

-- Query 7: Service Add-on Portfolio Depth & Customer Stickiness
SELECT NumServices, COUNT(*) AS customer_count, SUM(ChurnBinary) AS churned_count,
       ROUND(AVG(ChurnBinary) * 100.0, 2) AS churn_rate_pct
FROM customers GROUP BY NumServices ORDER BY NumServices ASC;

-- Query 8: Customer Lifecycle Cohort Churn Rates
SELECT TenureCohort, COUNT(*) AS accounts, SUM(ChurnBinary) AS lost_accounts,
       ROUND(AVG(ChurnBinary) * 100.0, 2) AS churn_rate_pct,
       ROUND(AVG(MonthlyCharges), 2) AS avg_monthly_bill
FROM customers GROUP BY TenureCohort ORDER BY MIN(tenure) ASC;

-- Query 9: Critical Defection Danger Zone Isolation
SELECT COUNT(*) AS danger_accounts, SUM(ChurnBinary) AS defected_accounts,
       ROUND(AVG(ChurnBinary) * 100.0, 2) AS danger_churn_rate,
       ROUND(SUM(MonthlyCharges), 2) AS monthly_revenue_at_risk
FROM customers
WHERE Contract = 'Month-to-month' AND InternetService = 'Fiber optic' AND PaymentMethod = 'Electronic check';

-- Query 10: High-Value Customer Defection Exposure (> $80.00/month)
SELECT Contract, COUNT(*) AS premium_customers, SUM(ChurnBinary) AS churned_premium,
       ROUND(AVG(ChurnBinary) * 100.0, 2) AS churn_rate_pct,
       ROUND(SUM(MonthlyCharges), 2) AS total_premium_mrr
FROM customers WHERE MonthlyCharges > 80.00 GROUP BY Contract;
```

APPENDIX II

MATHEMATICAL FORMULATIONS & STATISTICAL DERIVATIONS

1. Pearson's Chi-Square (Chi^2) Test of Independence

To test independence between two categorical variables with *r* rows and *c* columns, the test statistic Chi^2 evaluates discrepancy between observed cell counts (O_ij) and expected frequencies (E_ij):

Chi^2 = Sum Sum [(O_ij - E_ij)^2 / E_ij]

where expected frequencies under H0 are derived from marginal totals: **E_ij = (R_i * C_j) / N** with **d.f. = (r - 1)(c - 1)**.

2. Cramer's V Effect Size Coefficient

Because large samples (N = 7,043) inflate Chi-Square statistics, Cramer's V scales Chi^2 to bounded interval [0, 1]:

V = sqrt(Chi^2 / (N * min(r - 1, c - 1)))

Interpretation: V < 0.10: Negligible; 0.10 <= V < 0.20: Weak; 0.20 <= V < 0.35: Moderate; V >= 0.35: Strong association.

3. Welch's Two-Sample Independent T-Test

Welch's t-test evaluates whether two independent groups with unequal variances possess significantly different means:

t = (x_bar_1 - x_bar_2) / sqrt((s1^2 / N1) + (s2^2 / N2))

The degrees of freedom (df) are approximated via the Welch-Satterthwaite equation:

df = [(s1^2/N1 + s2^2/N2)^2] / [(s1^2/N1)^2/(N1 - 1) + (s2^2/N2)^2/(N2 - 1)].

4. Cohen's d Effect Size for Mean Differences

Cohen's d quantifies standardized distance between group means using pooled standard deviation (s_pooled):

d = (x_bar_1 - x_bar_2) / s_pooled

where **s_pooled = sqrt([(N1 - 1)*s1^2 + (N2 - 1)*s2^2] / (N1 + N2 - 2))**. Benchmarks: |d| ~ 0.2: Small; |d| ~ 0.5: Medium; |d| >= 0.8: Large effect size.

APPENDIX III

COMPLETE CLEANED DATA DICTIONARY & FEATURE SCHEMA

Attribute	Type	Domain Values	Analytical Description
customerID	String	7,043 unique alphanumeric IDs	Primary account identifier
gender	String	'Male', 'Female'	Customer biological sex
SeniorCitizen	Integer	0, 1	Senior citizen flag (1=Senior)
Partner	String	'Yes', 'No'	Domestic partner or spouse
Dependents	String	'Yes', 'No'	Financial dependents indicator
tenure	Integer	0 to 72 months	Subscription duration in months
PhoneService	String	'Yes', 'No'	Fixed telephone service
MultipleLines	String	'Yes', 'No', 'No phone'	Multiple phone lines flag
InternetService	String	'DSL', 'Fiber optic', 'No'	Core broadband technology
OnlineSecurity	String	'Yes', 'No', 'No internet'	Cybersecurity firewall add-on
OnlineBackup	String	'Yes', 'No', 'No internet'	Cloud backup service add-on
DeviceProtection	String	'Yes', 'No', 'No internet'	Hardware warranty protection
TechSupport	String	'Yes', 'No', 'No internet'	Priority technical IT assistance
StreamingTV	String	'Yes', 'No', 'No internet'	Television streaming package
StreamingMovies	String	'Yes', 'No', 'No internet'	Movie streaming subscription
Contract	String	'Month-to-month', 'One year', 'Two year'	Contract agreement term
PaperlessBilling	String	'Yes', 'No'	Digital invoice statement flag
PaymentMethod	String	Electronic check, Mailed check, Bank transfer, Credit card	Invoice settlement channel
MonthlyCharges	Float	$18.25 to $118.75	Monthly recurring billing fee
TotalCharges	Float	$0.00 to $8,684.80	Cumulative lifetime charges billed
Churn	String	'Yes', 'No'	Target customer defection flag
ChurnBinary	Integer	1, 0	Engineered binary churn flag
TenureCohort	String	'0-12m', '13-24m', '25-48m', '49-72m'	Lifecycle tenure interval segment
MonthlyChargesTier	String	'Low (<$35)', 'Medium', 'High (>$70)'	Spend level classification tier
NumServices	Integer	0 to 8	Count of active add-on services

APPENDIX IV

POWERPOINT PRESENTATION SLIDES OUTLINE (PART 1: SLIDES 1 TO 6)

As mandated by the SNIST Project Review Committee guidelines, a 12-slide executive presentation was prepared. The first six slides detail problem formulation, motivation, literature survey, and system architecture:

Slide 1: Title Slide • Customer Churn Analysis & BI Dashboard • G Mythili, K Ashiwith, D Niyati • Guide: Mr. Syed. Mazharuddin • SNIST CSE-DS	**Slide 2: Problem Statement** • High Telecom Churn: 26.54% attrition rate • $139k Monthly Recurring Revenue lost • Failure of heuristic rules & black-box ML
Slide 3: Motivation & Objectives • Retaining accounts is 5-7x cheaper than acquisition • Build an interpretable, transparent SQL BI stack • Deliver actionable executive retention pillars	**Slide 4: Literature Survey Summary** • Review of static rules vs ML classifiers • Identified lack of causal transparency in ML • Justification for pure Data Science & SQL
Slide 5: Research Gap & Scope • Existing models neglect billing channel friction • Failure to link churn predictions to ROI dollars • Absence of production executive dashboards	**Slide 6: Proposed System** • 4-tier architecture: Ingest -> SQLite -> Stats -> UI • 28 cleaned & engineered attributes • Chi-Square & Welch's T-Test hypothesis testing

APPENDIX IV (CONTINUED)

POWERPOINT PRESENTATION SLIDES OUTLINE (PART 2: SLIDES 7 TO 12)

This section documents slides 7 through 12, covering implementation methodology, statistical results, dashboard demo, and future roadmap:

Slide 7: System Architecture • End-to-end data flow diagram • SQLite with B-tree indexing for instant queries • Streamlit reactive state & caching pipeline	**Slide 8: Methodology & Engineering** • Whitespace imputation for TotalCharges (tenure=0) • 1.5xIQR boundary outlier validation • Parametric & non-parametric statistical inference
Slide 9: Implementation Modules • data_cleaning.py, sql_analysis.py • descriptive_statistics.py, eda.py • streamlit_app.py with 9 interactive pages	**Slide 10: Results & Analysis** • Month-to-month contracts drive 88.5% of churn • Electronic check friction drives 45.3% churn • Tech Support cuts Fiber Optic churn by 70%
Slide 11: Conclusion & Future Scope • 3% churn cut preserves $188k/yr recurring revenue • Future integration of real-time Kafka streams • Automated webhook SMS retention triggers	**Slide 12: Project Demo & Queries** • Live Streamlit BI Dashboard demonstration • Interactive SQL sandbox & ROI simulation • Faculty Committee Q&A session

APPENDIX V

SOFTWARE ARCHITECTURE & SOURCE CODE MANIFEST

In compliance with project report documentation standards for software development projects, the complete module directory and source code manifest for the Customer Churn Analysis and BI Dashboard codebase is cataloged below:

Source File Path	Core Module Class / Script	Lines of Code	Operational Responsibility & Functions
src/data_cleaning.py	TelcoDataCleaner	265 lines	Schema auditing, whitespace type coercion, 1.5xIQR outlier fences, feature store enrichment, SQLite database loading.
src/descriptive_statistics.py	TelcoStatsAnalyzer	310 lines	Parametric mean/std/skewness/kurtosis, Pearson/Spearman correlations, Chi-Square independence, Cramer's V, Welch's t-test.
src/sql_analysis.py	TelcoSQLAnalyzer	340 lines	SQLite relational database engine, B-Tree indexes, execution of 10 business SQL queries, financial retention ROI calculator.
src/eda.py	Plotly Visualization Engine	420 lines	Exact split-donut charts, contract MRR horizontal progress bars, tenure scatter plots, correlation heatmaps, static image exports.
app/streamlit_app.py	Streamlit BI Dashboard	1,280 lines	Full 9-page reactive SaaS analytics interface, dark theme CSS injection, popover filters, live cohort telemetry.
src/generate_pdf_report.py	Academic PDF Generator	1,950 lines	Compiles official 48-page academic internship report strictly conforming to SNIST formatting, double borders, Roman/Arabic numbering.

Table A.1: Complete Software Module & Source Code Manifest.

REFERENCES

[1] Berson, A., Smith, S., & Thearling, K., *Building Data Mining Applications for CRM*, McGraw-Hill, New York, 2000.

[2] Blattberg, R. C., Kim, B. D., & Neslin, S. A., *Database Marketing: Analyzing and Managing Customers*, Springer Science & Business Media, New York, 2008.

[3] Breiman, L., 'Random Forests', *Machine Learning Journal*, 45(1): 5-32, Oct. 2001.

[4] Chen, T., & Guestrin, C., 'XGBoost: A Scalable Tree Boosting System', in *Proc. of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pp. 785-794, Aug. 2016.

[5] Dal Pozzolo, A., Caelen, O., Johnson, R. A., & Bontempi, G., 'Calibrating Probability with Undersampling for Unbalanced Classification', *IEEE Symposium Series on Computational Intelligence*, pp. 159-166, Dec. 2015.

[6] Geron, A., *Hands-On Machine Learning with Scikit-Learn, Keras & TensorFlow*, 2nd Edition, O'Reilly Media, Sebastopol, CA, 2019.

[7] IBM Corporation, 'Telco Customer Churn Public Analytical Dataset', Cognos Analytics Sample Data, 2018. Available: *https://www.kaggle.com/datasets/blastchar/telco-customer-churn*.

[8] K. M. Chandy and L. Lamport, 'Distributed snapshots: Determining global states of distributed systems', *ACM Transactions on Computer Systems*, 3(1): 63-75, Feb. 1985.

[9] Lynch, A. B. (Ed.), *Technical Writing and Academic Documentation Guidelines*, Prentice Hall, London, 1972.

[10] McKinney, W., *Python for Data Analysis: Data Wrangling with Pandas, NumPy, and Jupyter*, 3rd Edition, O'Reilly Media, 2022.

[11] Plotly Technologies Inc., 'Collaborative Data Science and Interactive Visual Web Graphics', Montreal, QC, 2024. Available: *https://plotly.com/python/*.

[12] Reinartz, W., & Kumar, V., 'The Impact of Customer Relationship Characteristics on Profitable Lifetime Duration', *Journal of Marketing*, 67(1): 77-99, Jan. 2003.

[13] Rudin, C., 'Stop Explaining Black Box Machine Learning Models for High Stakes Decisions and Use Interpretable Models Instead', *Nature Machine Intelligence*, 1(5): 206-215, May 2019.

[14] Sreenidhi Institute of Science and Technology (SNIST), *Guidelines for the Preparation and Presentation of Seminar and Project Reports for B.Tech Academic Programmes*, Department of Data Science (CSE-DS), Hyderabad, Oct. 2026.

[15] Stallings, W., *Data and Computer Communications*, 5th Edition, Prentice Hall of India (PHI) Publication, New Delhi, 2003.

[16] Streamlit Inc., 'Streamlit: The Fastest Way to Build and Share Data Apps', Snowflake Inc., 2024. Available: *https://docs.streamlit.io/*.

[17] Tukey, J. W., *Exploratory Data Analysis*, Addison-Wesley Publishing Company, Reading, MA, 1977.

[18] Welch, B. L., 'The Generalization of 'Student's' Problem when Several Different Population Variances are Involved', *Biometrika*, 34(1/2): 28-35, Jan. 1947.

[19] Cramer, H., *Mathematical Methods of Statistics*, Princeton University Press, Princeton, NJ, 1946.

[20] SQLite Development Team, 'SQLite In-Process Relational Database Architecture and B-Tree Implementation', 2024. Available: *https://www.sqlite.org/arch.html*.